

Gradient Boosting from Scratch

Oblivious Trees, Ordered Boosting, and Posterior Sampling

For readers who know calculus but not gradient boosting

Abstract

This document builds modern gradient boosting from the ground up, assuming only undergraduate calculus (derivatives, chain rule, Taylor expansion). We derive every equation step by step, implement each component in pure NumPy, and map every line of code back to its mathematical origin. Starting from Vanilla GBM, we introduce five innovations in sequence — Newton Step, Oblivious Trees, Ordered Target Statistics, Ordered Boosting, and Posterior Sampling (SGLB) — each motivated by a concrete limitation of the previous approach. Source code: [github.com/\[repo\]](https://github.com/[repo]).

Notation

Symbol	Meaning
i	sample index ($i = 1, \dots, N$)
f	feature index ($f = 1, \dots, F$)
t	boosting iteration index ($t = 1, \dots, T$)
d	tree depth level ($d = 0, \dots, D - 1$)
$x_i \in \mathbb{R}^F$	feature vector of sample i
$x_{i,f}$	value of feature f for sample i
y_i	target value of sample i
F_t	ensemble model after t iterations
h_t	single tree fitted at iteration t
g_i^t	positive gradient $\partial L(y_i, F_i) / \partial F_i$ of sample i at iteration t
r_i^t	pseudo-residual $r_i = -g_i$ (negative gradient) of sample i at iteration t
h_i^t	Hessian of sample i at iteration t
τ	split threshold (a real-valued scalar)
(f^*, τ^*)	optimal feature and threshold found by split search
(f_d, τ_d)	shared split at depth d of an Oblivious Tree
R_j	leaf region j
γ_j	leaf value in region R_j
σ	random permutation of $\{1, \dots, N\}$
$\hat{x}_{i,f}$	Ordered TS encoding of feature f for sample i

1 Background: What Problem Are We Solving?

We have a dataset $\{(x_i, y_i)\}_{i=1}^N$ where $x_i \in \mathbb{R}^F$ are features and $y_i \in \mathbb{R}$ are targets. We want to find a function $F : \mathbb{R}^F \rightarrow \mathbb{R}$ that minimizes the average loss:

$$\mathcal{L}(F) = \sum_{i=1}^N L(y_i, F(x_i))$$

For regression we use **Mean Squared Error**: $L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$.

The factor $\frac{1}{2}$ is a convenience that cancels the 2 from differentiation, but does not affect the minimizer.

We are not looking for a single parametric function (like a linear model). Instead, we will build F as a *sum of many small decision trees*, each correcting the mistakes of the previous ones. This is the core idea of **boosting**.

2 Vanilla Gradient Boosting Machine

2.1 Gradient Descent in Function Space

You may be familiar with gradient descent for a parameter vector θ :

$$\theta \leftarrow \theta - \alpha \cdot \nabla_{\theta} \mathcal{L}$$

GBM does the same thing, but the “parameter” being updated is the *function* F itself. We treat $F(x_i)$, the prediction at each training point, as a variable we want to optimize.

The gradient of $\mathcal{L}$ with respect to $F(x_i)$ is simply the partial derivative treating $F(x_i)$ as a scalar:

$$\frac{\partial \mathcal{L}}{\partial F(x_i)} = \frac{\partial}{\partial F(x_i)} \sum_j L(y_j, F(x_j)) = \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)}$$

(All other terms vanish because they do not depend on $F(x_i)$.)

Let $g_i^t = \partial L(y_i, F_{t-1}(x_i)) / \partial F_{t-1}(x_i)$ be the **positive gradient** (first derivative of loss with respect to the prediction). The **pseudo-residual** is the negative gradient:

$$r_i^t = -g_i^t = -\left. \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right|_{F=F_{t-1}}$$

MSE case. $L(y_i, F_i) = \frac{1}{2}(y_i - F_i)^2$, so:

$$g_i^t = \frac{\partial L}{\partial F_i} = F_i - y_i \quad \Rightarrow \quad r_i^t = -g_i^t = y_i - F_{t-1}(x_i)$$

The pseudo-residual $r_i^t = y_i - F_{t-1}(x_i)$ is the residual (prediction error). Fitting trees to pseudo-residuals r_i^t is equivalent to taking a step in the negative gradient direction in function space.

Listing 1 Pseudo-residual computation (`vanilla_gbm.py`)

```
1 def _get_gradients(self, y, F_pred):
2     # r_i = y_i - F_{t-1}(x_i) (pseudo-residual = negative gradient)
3     # For MSE: g_i = F_i - y_i, so r_i = -g_i = y_i - F_i
4     return y - F_pred          # shape (N,)
```

2.2 Initialization

Before the first tree, we predict a constant F_0 for all samples. The optimal constant minimizes the total loss:

$$F_0 = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$$

For MSE, take the derivative with respect to γ and set it to zero:

$$\frac{d}{d\gamma} \sum_{i=1}^N \frac{1}{2} (y_i - \gamma)^2 = - \sum_{i=1}^N (y_i - \gamma) = 0 \quad \Rightarrow \quad F_0 = \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$$

Listing 2 Initialization (vanilla_gbm.py)

```
1 def _initialize(self, y):
2     # F_0 = mean(y): the optimal constant prediction under MSE.
3     # Derivation: argmin_gamma sum (y_i - gamma)^2
4     #             d/d(gamma) = -2 * sum(y_i - gamma) = 0
5     #             => gamma* = mean(y)
6     return np.mean(y)
```

2.3 Fitting a Tree to Pseudo-residuals

Recall what we are trying to do: take one step of gradient descent in function space. The ideal update would be to add the pseudo-residual $r_i^t = -g_i^t$ to the prediction at every training point x_i , moving each prediction in the direction that reduces loss. The problem is that this update is defined only at the N training points — it tells us nothing about how to predict at a new, unseen x .

This is why we fit a **decision tree** h_t to the pseudo-residuals r_i^t . The tree generalises the gradient step to the entire input space: instead of memorising r_i^t at each point, it finds regions of feature space where the pseudo-residuals are similar, and assigns a single constant to each region. Predicting with h_t at a new point x is equivalent to asking “which region does x fall into, and what was the average correction needed there during training?”

Formally, h_t partitions the feature space into J disjoint leaf regions $\{R_j\}_{j=1}^J$ and outputs a constant γ_j in each region:

$$h_t(x) = \sum_{j=1}^J \gamma_j \cdot \mathbf{1}[x \in R_j]$$

2.3.1 Optimal Leaf Value

Within leaf R_j , the optimal constant γ_j minimizes the MSE over the pseudo-residuals assigned to that leaf:

$$\gamma_j = \arg \min_{\gamma} \sum_{i \in R_j} (r_i - \gamma)^2$$

Differentiating and setting to zero (same as the F_0 derivation):

$$\gamma_j = \frac{1}{|R_j|} \sum_{i \in R_j} r_i = \bar{r}_{R_j}$$

The optimal leaf value is the *mean pseudo-residual* in that leaf.

Listing 3 Leaf value computation (`vanilla_gbm.py`)

```

1 def _make_leaf(self, r):
2     # gamma_j = mean(r) over leaf region R_j
3     # r_i = y_i - F_i (pseudo-residual); argmin_gamma sum(r_i - gamma)^2
4     node = Node()
5     node.value = np.mean(r)    # gamma_j = r-bar_{R_j}
6     return node

```

2.3.2 Split Gain Derivation

How do we choose *which* feature and threshold to split on? We want the split that most reduces the total MSE over pseudo-residuals.

A split on feature f at threshold τ creates:

$$R_1 = \{i : x_{i,f} \leq \tau\}, \quad R_2 = \{i : x_{i,f} > \tau\}$$

The MSE in a region R with optimal leaf value $\bar{r}_R$ is:

$$\text{MSE}(R) = \sum_{i \in R} (r_i - \bar{r}_R)^2$$

Expanding the square:

$$\text{MSE}(R) = \sum_{i \in R} r_i^2 - 2\bar{r}_R \sum_{i \in R} r_i + |R|\bar{r}_R^2 = \sum_{i \in R} r_i^2 - |R|\bar{r}_R^2$$

where we used $\bar{r}_R = \frac{1}{|R|} \sum_{i \in R} r_i$, so $|R|\bar{r}_R^2 = \frac{(\sum_{i \in R} r_i)^2}{|R|}$.

We define the **split gain** as the reduction in total MSE achieved by the split:

$$\text{Gain}(f, \tau) = \text{MSE}(\text{parent}) - [\text{MSE}(R_1) + \text{MSE}(R_2)]$$

We want to choose the split that *maximizes* this reduction. Substituting the expanded form of $\text{MSE}(R)$:

$$\text{Gain}(f, \tau) = \left[\sum_i r_i^2 - \frac{(\sum_i r_i)^2}{N} \right] - \left[\sum_{i \in R_1} r_i^2 - \frac{(\sum_{i \in R_1} r_i)^2}{|R_1|} \right] - \left[\sum_{i \in R_2} r_i^2 - \frac{(\sum_{i \in R_2} r_i)^2}{|R_2|} \right]$$

The $\sum_i r_i^2$ terms cancel (parent = $R_1 \cup R_2$, disjoint), leaving:

$$\text{Gain}(f, \tau) = \frac{\left(\sum_{i \in R_1} r_i \right)^2}{|R_1|} + \frac{\left(\sum_{i \in R_2} r_i \right)^2}{|R_2|} - \frac{\left(\sum_i r_i \right)^2}{N}$$

The third term is the **baseline score** (parent MSE contribution). It is subtracted so that $\text{Gain} > 0$ only when splitting actually helps.

Why subtract the baseline? Without it, any split would always look beneficial, because $\frac{A^2}{n_1} + \frac{B^2}{n_2} \geq \frac{(A+B)^2}{n_1+n_2}$ by the Cauchy–Schwarz inequality (equality holds only when $A/n_1 = B/n_2$, i.e. the split produces two perfectly balanced groups with identical means — which means the split does nothing useful). Subtracting the baseline makes $\text{Gain} = 0$ in that degenerate case, and $\text{Gain} > 0$ only when the split genuinely separates the pseudo-residuals. In code, `score_no_split = sum_t**2 / N` is this baseline, computed once and reused across all candidate splits.

Listing 4 Split search (vanilla_gbm.py)

```

1 def _find_split(self, X, g):
2     N, F = X.shape
3     sum_t = g.sum()
4     # Baseline: score if we do NOT split (one leaf for all)
5     # = (sum g)^2 / N
6     score_no_split = sum_t ** 2 / N
7
8     best_gain, best_feature, best_threshold = 0.0, None, None
9
10    for j in range(F):
11        order = np.argsort(X[:, j])
12        g_sorted = g[order]
13        x_sorted = X[order, j]
14        sum_L = 0.0
15
16        for i in range(1, N):
17            sum_L += g_sorted[i - 1]
18            sum_R = sum_t - sum_L
19            n_L, n_R = i, N - i
20
21            # skip if x values are equal (no real threshold here)
22            if x_sorted[i] == x_sorted[i - 1]:
23                continue
24
25            # Gain = (sum_L)^2/n_L + (sum_R)^2/n_R - baseline
26            gain = sum_L**2 / n_L + sum_R**2 / n_R - score_no_split
27
28            if gain > best_gain:
29                best_gain = gain
30                best_feature = j
31                # threshold = midpoint between adjacent values
32                best_threshold = (x_sorted[i - 1] + x_sorted[i]) / 2.0
33
34    return best_feature, best_threshold, best_gain

```

2.3.3 What is h_t exactly?

After the split search finds the best (f^*, τ^*) at each node, the tree h_t is fully defined by two things:

1. **The partition.** The recursive splits divide the feature space into J disjoint leaf regions $R_1, \dots, R_J$. Every training sample x_i falls into exactly one leaf.
2. **The leaf values γ_j .** Within each leaf, the tree outputs a single constant — the mean

pseudo-residual $\bar{g}_{R_j}$ derived above.

Putting these together:

$$h_t(x) = \sum_{j=1}^J \bar{r}_{R_j} \cdot \mathbf{1}[x \in R_j]$$

h_t is *not* trying to predict y directly. It is predicting the **current pseudo-residual** — how much the existing model F_{t-1} is wrong by, in each region of feature space. Adding $\eta \cdot h_t$ to F_{t-1} nudges the predictions toward y , one small step at a time.

2.4 Model Update

After fitting tree h_t , we update the model:

$$F_t(x) = F_{t-1}(x) + \eta \cdot h_t(x)$$

where $\eta \in (0, 1]$ is the **learning rate** (step size). The final prediction after T iterations is:

$$F_T(x) = \underbrace{\bar{y}}_{F_0} + \eta \sum_{t=1}^T h_t(x)$$

Listing 5 Training loop (vanilla_gbm.py)

```

1 def fit(self, X, y):
2     self.F0 = self._initialize(y)           # F_0 = mean(y)
3     F_pred = np.full(len(y), self.F0)      # F_0(x_i) for all i
4
5     for _ in range(self.n_estimators):
6         r = self._get_gradients(y, F_pred)   # r_i = y_i - F_{t-1}(x_i)
7         ↪ (pseudo-residual)
8
9         tree = DecisionTree(max_depth=self.max_depth)
10        tree.fit(X, r)                       # find t*, compute gamma_j = r-bar
11
12        update = tree.predict(X)             # h_t(x_i) for all i
13        F_pred += self.learning_rate * update # F_t = F_{t-1} + eta * h_t
14        self.trees.append(tree)

```

Notation clarification: h_t versus F_t . These two symbols are easy to confuse because they both represent “the model at step t ”, but they refer to very different things:

Symbol	What it is	Used for
h_t	The single tree trained at iteration t	Correcting the current residual
F_t	The cumulative ensemble up to iteration t	Making the final prediction

The relationship is $F_t = F_{t-1} + \eta \cdot h_t$, or equivalently:

$$F_T(x) = \underbrace{\bar{y}}_{h_0 \text{ (constant)}} + \eta h_1(x) + \eta h_2(x) + \cdots + \eta h_T(x)$$

Each h_t is a weak learner (shallow tree fitting current residuals). F_T is the strong learner built by summing them all. In code, `F_pred` tracks F_t and grows with each iteration, while `tree` (i.e. h_t) is freshly created each time.

Algorithm 1 Vanilla GBM (MSE)

Require: $\{(x_i, y_i)\}_{i=1}^N$, learning rate η , number of trees T

```

1:  $F_0 \leftarrow \bar{y}$  ▷ optimal constant: _initialize()
2: for  $t = 1, 2, \dots, T$  do
3:    $r_i^t \leftarrow y_i - F_{t-1}(x_i)$  for all  $i$  ▷ pseudo-residuals ( $= -g_i^t$ ): _get_gradients()
4:    $(f^*, \tau^*) \leftarrow \arg \max_{f, \tau} \text{Gain}(f, \tau)$  ▷ best split: _find_split()
5:    $\gamma_j \leftarrow \bar{r}_{R_j}$  for each leaf  $j$  ▷ leaf values: _make_leaf()
6:    $h_t(x) \leftarrow \sum_j \gamma_j \mathbf{1}[x \in R_j]$ 
7:    $F_t(x) \leftarrow F_{t-1}(x) + \eta \cdot h_t(x)$  ▷ update: fit() loop
8: end for
9: return  $F_T$ 
```

Step	Equation	Function
Initialization	$F_0 = \bar{y}$	<code>_initialize()</code>
Pseudo-residuals	$r_i^t = y_i - F_{t-1}(x_i) = -g_i^t$	<code>_get_gradients()</code>
Split search	$(f^*, \tau^*) = \arg \max_{f, \tau} \text{Gain}(f, \tau)$	<code>_find_split()</code>
Leaf values	$\gamma_j = \bar{r}_{R_j}$	<code>_make_leaf()</code>
Update	$F_t = F_{t-1} + \eta \cdot h_t$	<code>fit()</code> loop

3 Newton Step and General Loss

3.1 Second-Order Taylor Approximation

When computing the optimal leaf value γ_j , instead of minimising the exact loss, we use a **second-order Taylor approximation** around the current prediction $F_{t-1}(x_i)$:

$$L(y_i, F_{t-1}(x_i) + \gamma) \approx L_i + g_i \gamma + \frac{1}{2} h_i \gamma^2$$

where:

$$g_i = \left. \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right|_{F_{t-1}} \quad (\text{first derivative, gradient})$$

$$h_i = \left. \frac{\partial^2 L(y_i, F(x_i))}{\partial F(x_i)^2} \right|_{F_{t-1}} \quad (\text{second derivative, Hessian})$$

Minimising the approximation over γ :

$$\frac{d}{d\gamma} (L_i + g_i \gamma + \frac{1}{2} h_i \gamma^2) = g_i + h_i \gamma = 0 \quad \Rightarrow \quad \gamma = -\frac{g_i}{h_i}$$

Summing over all samples in leaf R_j :

$$\gamma_j = -\frac{\sum_{i \in R_j} g_i}{\sum_{i \in R_j} h_i}$$

MSE verification. $g_i = F_i - y_i$, $h_i = 1$. $\gamma_j = -\frac{\sum g_i}{|R_j|} = \frac{\sum (y_i - F_i)}{|R_j|} = \bar{r}_{R_j}$ (mean pseudo-residual). Consistent with Vanilla GBM.

3.2 Loss Interface

We implement a unified interface so any loss function can be plugged in:

Loss	g_i	h_i	F_0	Notes
MSE ($\frac{1}{2}(y - F)^2$)	$F_i - y_i$	1	$\bar{y}$	regression
LogLoss ($-y \log p - (1 - y) \log(1 - p)$)	$p_i - y_i$	$p_i(1 - p_i)$	$\log \frac{\bar{y}}{1 - \bar{y}}$	binary classification

where $p_i = (1 + e^{-F_i})^{-1}$ is the sigmoid function.

Listing 6 Loss functions (loss.py)

```

1 class MSE:
2     def gradient(self, y, F):
3         return F - y                                # g_i = F_i - y_i = dL/dF
4
5     def hessian(self, y, F):
6         return np.ones_like(y)                      # h_i = d^2L/dF^2 = 1
7
8     def leaf_value(self, g, h):
9         return -g.sum() / h.sum()                   # gamma_j = -sum(g)/sum(h)
10
11     def init_prediction(self, y):
12         return np.mean(y)                           # F_0 = mean(y)
13
14 class LogLoss:
15     def gradient(self, y, F):
16         p = 1.0 / (1.0 + np.exp(-F))
17         return p - y                                # g_i = p_i - y_i
18
19     def hessian(self, y, F):
20         p = 1.0 / (1.0 + np.exp(-F))
21         return p * (1.0 - p)                        # h_i = p_i(1 - p_i)
22
23     def leaf_value(self, g, h):
24         return -g.sum() / (h.sum() + 1e-9)
25
26     def init_prediction(self, y):
27         p = np.clip(np.mean(y), 1e-7, 1 - 1e-7)
28         return np.log(p / (1.0 - p))                # F_0 = log-odds

```

4 Oblivious Trees

4.1 Motivation: Why Change the Tree Structure?

A standard decision tree chooses a *different* feature and threshold at each node. This gives high flexibility but also high variance — the tree can overfit by memorising specific regions of the training data.

CatBoost replaces standard trees with **Oblivious Trees**: trees in which every node at the same depth uses the *same* feature and threshold.

Think of an Oblivious Tree as a sequence of D yes/no questions applied to every sample:

1. Is `age > 35`?

2. Is `income > 50000`?
3. Is `credit_score > 700`?

The answers (0 or 1) for each sample form a D -bit binary number that directly identifies a leaf — no traversal needed.

Figure 1 compares the two structures at depth 3.




Figure 1: Standard vs Oblivious tree (depth 3, 8 leaves each). (a) Standard tree: the same feature `score` appears at depth 2, but each node uses a *different* threshold (600, 750, 680, 720). (b) Oblivious tree: every node at the same depth shares an *identical* condition (boxed). Three shared conditions yield a 3-bit index that directly selects one of $2^3 = 8$ leaves — no path traversal needed.

Both trees in Figure 1 have the same depth and the same number of leaves. The visual difference reflects a genuine difference in **expressiveness**: a standard tree of depth 3 can use up to $2^3 - 1 = 7$ independent split conditions, while an Oblivious Tree uses only 3 (one per depth level). This lower flexibility is a deliberate tradeoff — it acts as structural regularization, making each tree less prone to memorising noise. CatBoost compensates by building more trees in the ensemble rather than making each tree more complex.

4.2 Structure

An Oblivious Tree of depth D is defined by D pairs:

$$\text{splits} = [(f_0, \tau_0), (f_1, \tau_1), \dots, (f_{D-1}, \tau_{D-1})]$$

and 2^D leaf values $\{\gamma_j\}_{j=0}^{2^D-1}$.

4.3 Split Search

At depth d , all 2^d existing nodes share the same (f_d, τ_d) . This shared constraint changes how we search for the best split.

In a standard tree, we find the best split *independently* for each node: each node looks only at its own samples and picks whatever (f, τ) maximizes its local gain. Two sibling nodes can therefore choose completely different features and thresholds.

In an Oblivious Tree, one (f, τ) must serve *all* 2^d nodes simultaneously. The natural criterion is to pick the (f, τ) that maximizes the *total* gain summed across every node at this depth:

$$\text{Gain}(f, \tau) = \sum_{n=1}^{2^d} \left[\frac{\left(\sum_{i \in L_n} g_i \right)^2}{\sum_{i \in L_n} h_i} + \frac{\left(\sum_{i \in R_n} g_i \right)^2}{\sum_{i \in R_n} h_i} - \frac{\left(\sum_{i \in n} g_i \right)^2}{\sum_{i \in n} h_i} \right]$$

Each bracketed term is the gain from applying (f, τ) to one node n — identical in form to the Vanilla GBM gain. The $\sum_{n=1}^{2^d}$ simply adds these up across all 2^d nodes. A split that helps many nodes a little beats one that helps one node a lot but hurts the others.

Here h_i is the Hessian (second derivative of loss) — we will derive this in Section 3. For MSE, $h_i = 1$ and each bracketed term reduces exactly to the Vanilla GBM gain.

Listing 7 Oblivious split search (tree.py)

```
1 def _find_split(self, X, g, h, node_indices):
2     # node_indices: list of arrays, one per current node,
3     #               each containing the sample indices in that node.
4     _, F = X.shape
5     best_gain, best_f, best_t = 0.0, 0, 0.0
6
7     for j in range(F):                # try every feature
8         # Collect all x-values across all nodes to find thresholds
9         all_x = np.concatenate([X[idx, j] for idx in node_indices])
10        unique_x = np.unique(all_x)
11        if len(unique_x) < 2:
12            continue
13        thresholds = (unique_x[:-1] + unique_x[1:]) / 2.0
14
15        for thresh in thresholds:
16            gain_total = 0.0
17            valid = True
18
19            for idx in node_indices:    # SUM over all nodes at this depth
20                mask = X[idx, j] <= thresh
21                n_L, n_R = mask.sum(), (~mask).sum()
22                if n_L < self.min_samples_leaf or n_R < self.min_samples_leaf:
23                    valid = False; break
24
25                gL = g[idx][mask].sum(); gR = g[idx][~mask].sum()
26                hL = h[idx][mask].sum(); hR = h[idx][~mask].sum()
27                gt = g[idx].sum(); ht = h[idx].sum()
28
29                # Newton gain: (sum_g)^2 / sum_h for each side
30                gain_total += (gL**2/(hL+1e-9) + gR**2/(hR+1e-9)
31                             - gt**2/(ht+1e-9))
32
33            if valid and gain_total > best_gain:
34                best_gain, best_f, best_t = gain_total, j, thresh
35
36        return best_f, best_t
```

4.4 Leaf Value

After finding all D splits, samples are assigned to 2^D leaves. The leaf value uses the Newton step (Section 3):

$$\gamma_j = -\frac{\sum_{i \in R_j} g_i}{\sum_{i \in R_j} h_i}$$

For MSE ($h_i = 1$): $\gamma_j = -\frac{\sum_{i \in R_j} g_i}{|R_j|} = \bar{g}_{R_j}$, identical to Vanilla GBM.

Listing 8 Leaf value assignment (`tree.py`)

```
1 n_leaves      = 2 ** self.max_depth
2 self.leaf_values = np.zeros(n_leaves)
3 for k, idx in enumerate(node_indices):
4     if len(idx) > 0:
5         # Newton step: gamma_j = -sum(g) / sum(h)
6         # For MSE h=1, this equals mean(y - F) = r-bar (mean pseudo-residual)
7         self.leaf_values[k] = -g[idx].sum() / (h[idx].sum() + 1e-9)
```

4.5 Bit-Index Prediction

After training, the tree is fully described by two things: D split conditions $(f_0, \tau_0), \dots, (f_{D-1}, \tau_{D-1})$, and an array of 2^D leaf values $\gamma_0, \dots, \gamma_{2^D-1}$.

To **predict** for a new sample x , we need to find which leaf it belongs to, then return that leaf's value γ_j . In a standard tree, this means traversing the tree from the root: at each node, check that node's condition and go left or right, then repeat at the child node, until we reach a leaf. This takes $O(D)$ steps, and each step depends on the previous one — the path is different for every sample.

In an Oblivious Tree, the symmetric structure enables a shortcut. Because every node at depth d tests the *same* condition (f_d, τ_d) , we can evaluate all D conditions *independently* and simultaneously. Each condition produces a single bit $b_d \in \{0, 1\}$:

$$b_d = \mathbf{1}[x_{f_d} > \tau_d]$$

The D bits together uniquely identify the leaf — they are exactly the sequence of left/right decisions the sample would make at each depth. Concatenating them as a binary number gives the leaf index directly:

$$\text{index}(x) = \sum_{d=0}^{D-1} b_d \cdot 2^{D-1-d}$$

The prediction is then a single array lookup: $h_t(x) = \gamma_{\text{index}(x)}$. No traversal, no branching — just D comparisons and one table lookup.

This is computed incrementally using left-shift accumulation (`index = index * 2 + bit`), which builds the binary integer one bit at a time from most significant to least significant.

Example. $D = 3$, splits evaluate to $(b_0, b_1, b_2) = (1, 0, 1)$:

$$\text{index} = 0 \xrightarrow{b_0=1} 0 \cdot 2 + 1 = 1 \xrightarrow{b_1=0} 1 \cdot 2 + 0 = 2 \xrightarrow{b_2=1} 2 \cdot 2 + 1 = 5$$

This sample lands in leaf 5 (out of $2^3 = 8$ leaves), and its prediction is γ_5 .

The same bit-index calculation appears in two places in the code. `predict()` uses it to return leaf values at inference time. `get_leaf_indices()` uses it during training to tell Ordered Boosting which leaf each sample falls into (needed for the per-leaf cumulative statistics in Section 6).

Listing 9 Prediction via bit-index (`tree.py`)

```
1 def predict(self, X):
2     indices = np.zeros(len(X), dtype=int) # start: index = 0 for all samples
3
4     for feature, threshold in self.splits:
5         # bit_d = 1 if x[feature] > threshold, else 0
6         bit = (X[:, feature] > threshold).astype(int)
7         # Left-shift and append bit: index = index * 2 + bit
8         # This builds the binary number d bit at a time.
9         indices = indices * 2 + bit
10
11     return self.leaf_values[indices] # O(1) lookup: no traversal
12
13 def get_leaf_indices(self, X):
14     # Same logic, but return the integer index instead of the leaf value.
15     # Used by Ordered Boosting to track which leaf each sample falls into.
16     indices = np.zeros(len(X), dtype=int)
17     for feature, threshold in self.splits:
18         indices = indices * 2 + (X[:, feature] > threshold).astype(int)
19     return indices
```

A standard decision tree traverses a path from root to leaf: $O(\text{depth})$ conditional checks, each depending on the previous. An Oblivious Tree evaluates D independent conditions and combines them with bit arithmetic — all D conditions can be evaluated *in parallel* across all samples, making it highly GPU-friendly.

It may seem counterintuitive that more leaves (2^D vs a shallower standard tree) leads to *faster* prediction. The key insight is that the speed gain does not come from having fewer leaves, but from *eliminating tree traversal*. In a standard tree, each sample follows a different path through the nodes, requiring sequential conditional branching. In an Oblivious Tree, every sample evaluates the same D conditions — so the entire dataset can be processed with D vectorised comparisons followed by a single array lookup. The 2^D leaf table is precomputed at training time and accessed in $O(1)$ at prediction time, regardless of D .

5 Ordered Target Statistics

5.1 The Target Leakage Problem

Many datasets contain categorical features (city, product type, user segment). A natural way to encode them is **target statistics**: replace each category value with the mean target among samples of that category:

$$\hat{x}_{k,i} = \frac{\sum_{j=1}^N \mathbf{1}[x_{j,i} = x_{k,i}] \cdot y_j}{\sum_{j=1}^N \mathbf{1}[x_{j,i} = x_{k,i}]}$$

The problem: when computing $\hat{x}_{k,i}$ for sample k , the sum includes $j = k$, i.e. the sample's own target y_k is used to encode its own feature.

Concrete example. Suppose category “City A” has 3 samples with targets $y = 10, 20, 30$. The encoding for the first sample ($y_1 = 10$) is:

$$\hat{x}_1 = \frac{10 + 20 + 30}{3} = 20$$

At training time the model sees feature value 20 and target 10 — the encoding is contaminated by its own label, making the model overconfident during training but degrading at test time when the sample is absent from the denominator.

5.2 Ordered TS: The Fix

Draw a random permutation σ of $\{1, \dots, N\}$. For sample k at position $r = \sigma^{-1}(k)$ in the permutation, use *only* the $r - 1$ samples that come *before* k in σ :

$$\hat{x}_{k,i} = \frac{\sum_{\substack{j : \sigma^{-1}(j) < r \\ x_{j,i} = x_{k,i}}} y_j + a \cdot p}{\sum_{\substack{j : \sigma^{-1}(j) < r \\ x_{j,i} = x_{k,i}}} 1 + a}$$

- $p = \bar{y}$: the global target mean (**prior**).
- $a > 0$: the **smoothing strength**. When few prior samples share the same category as k , the estimate defaults toward the global mean, preventing noisy estimates from rare categories.
- The strict inequality $\sigma^{-1}(j) < r$ ensures k itself is *never* included. Leakage is eliminated by construction.

Critical implementation detail: the cumulative accumulators must be updated *after* computing $\hat{x}_{k,i}$, not before.

Listing 10 Ordered TS transform (`encoder.py`)

```
1 def transform(self, X, y, sigma):
2     N = len(y)
3     X_enc = X.astype(float).copy()
4     p, a = self.prior_, self.prior_weight    # prior = mean(y), smoothing a
5
6     for col in self.cat_features:
7         cat_sum = defaultdict(float)    # sum of y_j for each category value
8         cat_cnt = defaultdict(int)      # count of preceding samples
9
10        for r in range(N):
11            k = sigma[r]                # process samples in permutation order
12            cat_val = int(X[k, col])
13
14            # --- STEP 1: compute encoding using ONLY preceding samples ---
15            cnt = cat_cnt[cat_val]      # how many j with sigma^{-1}(j) < r
16            s = cat_sum[cat_val]        # sum of y_j for those j
17
18            # x-hat_{k,col} = (sum_y_j + a*p) / (cnt + a)
19            X_enc[k, col] = (s + p * a) / (cnt + a)
20
21            # --- STEP 2: update accumulators AFTER encoding (y_k excluded) ---
22            cat_sum[cat_val] += y[k]
23            cat_cnt[cat_val] += 1
24
25        return X_enc
```

Think of the permutation as a timeline. Sample k arrives at time r . Its categorical encoding is computed from the history up to time $r - 1$. Its own outcome y_k is unknown at that moment, so it is excluded. This mirrors the real-world situation: at test time, y_k is unknown.

6 Ordered Boosting

Note. This section presents an educational approximation of the CatBoost-style ordered boosting idea. The implementation uses a permutation-based prefix to reduce self-influence (prediction shift), but it is not a complete reproduction of CatBoost’s production algorithm, which uses multiple permutations, separate model sets, and additional engineering. Treat the results as an illustration, not a benchmark.

6.1 Gradient Leakage in Standard GBM

In Vanilla GBM, the positive gradient of sample k at iteration t is:

$$g_k^t = F_{t-1}(x_k) - y_k$$

where F_{t-1} was trained on *all* N samples including k . Equivalently, the pseudo-residual is $r_k^t = y_k - F_{t-1}(x_k)$.

Because F_{t-1} has already fitted sample k , it predicts closer to y_k than it would for an unseen sample. This causes $|g_k^t|$ to be **underestimated** (i.e. the pseudo-residual r_k^t is closer to zero than it should be).

Numerical example. Suppose $y_k = 1.0$.

- F_{t-1} trained *including* k : $F_{t-1}(x_k) = 0.8 \Rightarrow r_k = 0.2$ (biased small)

- F_{t-1} trained *excluding* k : $F_{t-1}(x_k) = 0.5 \Rightarrow r_k = 0.5$ (less biased)

The biased pseudo-residual tells the next tree that sample k only needs a correction of 0.2, when the true correction needed is closer to 0.5. The model systematically under-corrects its mistakes on training data, which is a form of overfitting.

6.2 Ordered Boosting: The Fix

Draw a random permutation σ . For sample k at position r in σ , compute the positive gradient against a model M_{r-1}^t that was built using only the $r - 1$ samples preceding k in σ :

$$g_{\sigma(r)}^t = M_{r-1}^t(x_{\sigma(r)}) - y_{\sigma(r)}$$

Because M_{r-1}^t has never seen $x_{\sigma(r)}$, the gradient estimate has reduced self-influence compared to Vanilla GBM. (This is an approximation; the full CatBoost algorithm uses additional mechanisms to control this bias.)

6.3 Two Models: M_{r-1}^t versus F_{t-1}

Before deriving the per-leaf approximation, it is worth clarifying why we need *two* different models in Ordered Boosting, and what each one does. This is a common point of confusion.

Symbol	Purpose	Updated
F_{t-1}	Final ensemble for prediction	Once per iteration (after tree fit)
M_{r-1}^t	Gradient estimation only	Once per sample (within iteration)

F_{t-1} is the same model used in Vanilla GBM — the full ensemble accumulated over the previous $t - 1$ iterations. It is used at the *end* of iteration t to update the prediction: $F_t = F_{t-1} + \eta \cdot h_t$.

M_{r-1}^t is an auxiliary model used *within* iteration t to compute prefix-corrected gradients. It is built from samples $\sigma(1), \dots, \sigma(r - 1)$ — those that appear before sample k in the current permutation. Because M_{r-1}^t has never seen x_k , using it to estimate g_k^t reduces the self-influence that afflicts Vanilla GBM.

After all gradients $\{g_k^t\}$ are collected, a new tree h_t is fitted to them and F_t is updated — exactly as in Vanilla GBM. The role of M_{r-1}^t ends there; it is an internal tool for reduced-bias estimation, not part of the final model.

6.4 Per-Leaf Cumulative Statistics

Building a separate model for each sample is intractable (N models per iteration). We approximate M_{r-1}^t using the *previous iteration's tree* to assign leaf memberships, then accumulate residuals within each leaf in σ -order:

$$M_{r-1}^t(x_k) \approx F_{t-1}(x_k) + \frac{\sum_{\substack{s < r \\ x_{\sigma(s)} \in R_j}} (y_{\sigma(s)} - F_{t-1}(x_{\sigma(s)}))}{|\{s < r : x_{\sigma(s)} \in R_j\}|}$$

where R_j is the leaf of the previous tree containing x_k .

In code, $\text{base}_k = F_{t-1}(x_k) - y_k = g_k^t$ (the vanilla positive gradient), and **correction** is the mean base of previous samples in the same leaf. It is worth verifying that $\mathbf{g}[\mathbf{k}] = \mathbf{base}[\mathbf{k}] - \mathbf{correction}$ is indeed approximately equal to $M_{r-1}^t(x_k) - y_k$:

$$\begin{aligned}
g_k^{\text{ordered}} &= \text{base}_k - \text{correction}_k \\
&= (F_{t-1}(x_k) - y_k) - \underbrace{(F_{t-1}(x_k) - M_{r-1}^t(x_k))}_{\text{correction}} \\
&\approx M_{r-1}^t(x_k) - y_k
\end{aligned}$$

The correction term $F_{t-1}(x_k) - M_{r-1}^t(x_k)$ is approximated by the per-leaf running mean of preceding gradients: it measures how much the leaf-level gradients of preceding samples shift the model's prediction beyond F_{t-1} .

Listing 11 Ordered gradient computation (boosting.py)

```

1 class OrderedBoosting:
2     def compute_gradients(self, y, F_pred, loss, lr,
3                           prev_tree=None, sigma=None, X_enc=None):
4         N = len(y)
5         base = loss.gradient(y, F_pred)      # base_k = g_k = F_{t-1}(x_k) - y_k
6         ↪ (positive gradient)
7
8         if prev_tree is not None and X_enc is not None:
9             leaf_ids = prev_tree.get_leaf_indices(X_enc)
10            n_leaves = prev_tree.n_leaves
11        else:
12            leaf_ids = np.zeros(N, dtype=int)  # first iteration: all leaf 0
13            n_leaves = 1
14
15        leaf_sum = np.zeros(n_leaves)
16        leaf_cnt = np.zeros(n_leaves, dtype=int)
17        g = np.empty(N)
18
19        for r in range(N):
20            k = sigma[r]
21            li = leaf_ids[k]
22            corr = leaf_sum[li] / leaf_cnt[li] if leaf_cnt[li] > 0 else 0.0
23            g[k] = base[k] - corr              # g_k = base_k - correction
24
25            # Update AFTER computing g[k] (k excluded from own correction)
26            leaf_sum[li] += base[k]
27            leaf_cnt[li] += 1
28
29        h = loss.hessian(y, F_pred)
30        return g, h

```

This is the same idea as Ordered TS, applied to gradient estimation. In Ordered TS, sample k 's feature encoding excludes its own y_k . In Ordered Boosting, sample k 's gradient excludes its own contribution to the model. Both prevent a sample from “seeing” its own label.

Two nested loops, two different roles. It is easy to lose track of what the two σ -order loops (Steps 2 and 3) are doing. They share the same permutation σ but serve completely different purposes:

Algorithm 2 CatBoost-style ordered boosting with Oblivious Trees (educational approximation)

Require: $\{(x_i, y_i)\}_{i=1}^N$, η , T , depth D , smoothing a

```

1:  $F_0(x) \leftarrow \bar{y}$ ;  $h_0 \leftarrow \text{None}$ 
2: for  $t = 1, 2, \dots, T$  do
3:   // Step 1: fresh permutation
4:    $\sigma \leftarrow \text{RandomPermutation}(N)$ 
5:   // Step 2: Ordered TS — encode categorical features
6:    $S_c \leftarrow 0$ ,  $C_c \leftarrow 0$  for every category  $c$ 
7:   for  $r = 1, \dots, N$  do
8:      $k \leftarrow \sigma(r)$ 
9:     for each categorical feature  $i$  do
10:       $c \leftarrow x_{k,i}$ ;  $\hat{x}_{k,i} \leftarrow (S_c + a\bar{y}) / (C_c + a)$  ▷ only preceding samples;  $k$  excluded
11:       $S_c += y_k$ ;  $C_c += 1$  ▷ update AFTER encoding
12:    end for
13:  end for
14:   $\tilde{x}_i \leftarrow$  encoded features of sample  $i$ 
15:  // Step 3: Ordered Boosting — prefix-corrected gradients
16:   $j_i \leftarrow$  leaf of  $h_{t-1}$  for  $\tilde{x}_i$  ▷  $h_0 = \text{None}$ : all  $j_i \leftarrow 0$ 
17:   $A_j \leftarrow 0$ ,  $N_j \leftarrow 0$  for  $j = 0, \dots, 2^D - 1$ 
18:  for  $r = 1, \dots, N$  do
19:     $k \leftarrow \sigma(r)$ ;  $j \leftarrow j_k$ 
20:     $b_k \leftarrow \partial L / \partial F(y_k, F_{t-1}(\tilde{x}_k))$ ;  $\delta \leftarrow A_j / N_j$  if  $N_j > 0$  else 0
21:     $g_k^t \leftarrow b_k - \delta$  ▷ prefix-corrected:  $\approx \partial L / \partial F(y_k, M_{r-1}^t(\tilde{x}_k))$ 
22:     $A_j += b_k$ ;  $N_j += 1$  ▷ update AFTER;  $k$  excluded from own correction
23:  end for
24:   $h_i^t \leftarrow \partial^2 L / \partial F^2(y_i, F_{t-1}(\tilde{x}_i))$  ▷ Hessian; = 1 for MSE
25:  // Step 4: fit one Oblivious Tree
26:  All  $N$  samples in root; splits  $\leftarrow []$ 
27:  for  $d = 1, \dots, D$  do
28:     $(f^*, \tau^*) \leftarrow \arg \max_{f, \tau} \sum_n [\frac{G_{L_n}^2}{H_{L_n}} + \frac{G_{R_n}^2}{H_{R_n}} - \frac{G_n^2}{H_n}]$  ▷ one split shared across ALL  $2^{d-1}$  nodes
29:    splits.append( $f^*, \tau^*$ ); split every node on  $\tilde{x}_{f^*}$  vs  $\tau^*$ 
30:  end for
31:   $\gamma_j \leftarrow -G_j / H_j$  for each leaf  $j$  ▷  $G_j = \sum_{i \in j} g_i^t$ ,  $H_j = \sum_{i \in j} h_i^t$ 
32:   $h_t(x) \leftarrow \gamma_{\text{idx}(x)}$ ,  $\text{idx}(x) = \sum_{d=0}^{D-1} \mathbf{1}[\tilde{x}_{f_d} > \tau_d] \cdot 2^{D-1-d}$  ▷ bit-index
33:  // Step 5: update ensemble
34:   $F_t(x) \leftarrow F_{t-1}(x) + \eta \cdot h_t(x)$ ;  $h_{t-1} \leftarrow h_t$ 
35: end for
36: return  $F_T$ 

```

Loop	What it computes	Accumulates
Step 2 (Ordered TS)	$\hat{x}_{k,i}$: categorical encoding of features	S_c, C_c per category
Step 3 (Ordered Boosting)	g_k^t : prefix-corrected gradient	A_j, N_j per leaf

Both loops follow the same invariant: *compute first, update after*. This single rule is what prevents sample k from appearing in its own encoding (Step 2) or its own gradient correction (Step 3).

6.5 When Does Ordered Boosting Actually Help?

Ordered TS and Ordered Boosting solve *different* leakage problems, and it is important to keep them separate.

Ordered TS is specific to categorical features. When a category's encoding uses the sample's own y_k , the encoded feature becomes a direct function of the target — a strong and direct form of leakage. Without categorical features, Ordered TS is simply not needed.

Ordered Boosting addresses gradient leakage, which exists regardless of feature type. Even with purely numerical features, F_{t-1} was trained on all N samples including k , so $F_{t-1}(x_k)$ is biased toward y_k , causing g_k^t to be underestimated. This is a real problem.

However, the *magnitude* of the two leakages is very different. In Ordered TS, sample k 's own y_k appears directly in the numerator of its feature encoding — a first-order effect. In Ordered Boosting with numerical features, the bias in $F_{t-1}(x_k)$ is a weaker, indirect effect: the model has seen x_k during training, but shares it with $N - 1$ other samples, so the per-sample influence is $O(1/N)$.

In practice, this means:

- **Numerical features only, large dataset** — Ordered Boosting provides marginal benefit. Plain mode is faster and often equally accurate.
- **Categorical features present** — Ordered TS eliminates a strong leakage source, and Ordered Boosting should be used alongside it (the paper requires $\sigma_{\text{cat}} = \sigma_{\text{boost}}$ to guarantee full leakage removal).
- **Small dataset, numerical only** — Ordered Boosting can still reduce overfitting, since the per-sample influence of $O(1/N)$ is larger when N is small.
- **External walk-forward validation already applied** — If the training pipeline already prevents look-ahead bias at the split level (e.g. time-series walk-forward retraining), Ordered Boosting adds little on top.

The CatBoost documentation reflects this: Ordered mode is recommended for small datasets, while Plain mode is the default for large ones. For datasets with only numerical features and an existing temporal validation structure, Plain mode with other regularization (e.g. posterior sampling, Section 7) is often the better choice.

7 Posterior Sampling via Stochastic Gradient Langevin Boosting

7.1 Motivation: From Point Estimate to Distribution

Vanilla GBM finds a single model F^* that minimizes the loss:

$$F^* = \arg \min_F \sum_{i=1}^N L(y_i, F(x_i))$$

This is a point estimate. It tells us *what* to predict, but not *how confident* we should be. Three problems remain unaddressed:

- **Local minima.** When the loss landscape is non-convex, F^* may be a local rather than global minimum. Standard GBM has no mechanism to escape.
- **Overfitting.** Gradient boosting aggressively fits the training data. Without explicit regularization, F^* memorizes noise. Previous chapters addressed this structurally (Oblivious Trees, Ordered Boosting); SGLB adds a complementary stochastic regularization: noise injection prevents the model from converging too precisely to any single point, acting similarly to dropout in neural networks.
- **Uncertainty.** For a new sample x , how much should we trust the prediction $F^*(x)$? A point estimate gives no answer.

All three are addressed by replacing the point estimate with a **posterior distribution** over functions $p(F|\mathcal{D})$.

7.2 Langevin Dynamics

Langevin dynamics adds noise to gradient descent:

$$\theta_{t+1} = \theta_t - \eta \nabla L(\theta_t) + \sqrt{\frac{2\eta}{\beta}} \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(0, I)$$

The noise term $\sqrt{2\eta/\beta} \epsilon_t$ allows the algorithm to escape local minima — it performs a gradient-biased random walk rather than pure gradient descent.

The key theoretical result (via the Fokker-Planck equation) is that this Markov chain converges in distribution to:

$$p(\theta) \propto \exp(-\beta L(\theta))$$

This is a **Boltzmann distribution**: states with low loss have high probability, states with high loss have low probability. β is the *inverse temperature* — large β concentrates the distribution near the minimum (approaching MAP), small β spreads it broadly.

7.3 Connection to the Posterior

If we define loss as negative log-likelihood, $L(F) = -\log p(\mathcal{D}|F)$, and assume a uniform prior, then:

$$p(F|\mathcal{D}) \propto p(\mathcal{D}|F) \cdot p(F) = \exp(-L(F))$$

The stationary distribution of Langevin dynamics is exactly the **Bayesian posterior** over F . Gradient descent finds the posterior mode (MAP); Langevin dynamics *samples* from the full posterior.

7.4 SGLB: Applying Langevin Dynamics to Boosting

Applying Langevin dynamics to function space gives **Stochastic Gradient Langevin Boosting** (SGLB, Ustimenko & Prokhorenkova, 2021). Three changes from Vanilla GBM:

1. **Noise injection.** Add Langevin noise to positive gradients:

$$\tilde{g}_k^t = g_k^t + \epsilon_k, \quad \epsilon_k \sim \mathcal{N}\left(0, \frac{2\eta}{\beta}\right)$$

2. **Noisy tree fit.** Fit the tree to $\tilde{g}^t$ instead of g^t :

$$h_t = \text{FitTree}(X, \tilde{g}^t)$$

3. **Shrinkage update.** Apply L2-style weight decay to F :

$$F_t = (1 - \alpha\eta) F_{t-1} + \eta \cdot h_t(x)$$

The shrinkage term $(1 - \alpha\eta)$ ensures the stationary distribution is normalizable for any loss function:

$$p(F) \propto \exp\left(-\beta \left(L(F) + \frac{\alpha}{2} \|F\|^2\right)\right)$$

Why shrinkage is necessary. Without shrinkage, the stationary distribution $\exp(-\beta L(F))$ may not be normalizable — for example, if $L(F) \rightarrow 0$ as $F \rightarrow -\infty$ (as with exponential-family losses), the integral $\int \exp(-\beta L(F)) dF$ diverges. The $\frac{\alpha}{2}\|F\|^2$ term guarantees integrability regardless of the loss function.

Listing 12 SGLB gradient and update (boosting.py)

```

1 class SGLB:
2     def __init__(self, alpha=0.01, beta=100.0, rng=None):
3         self.alpha = alpha          # shrinkage strength
4         self.beta  = beta            # inverse temperature
5         self.rng   = rng or np.random.default_rng()
6
7     def compute_gradients(self, y, F_pred, loss, lr,
8                           prev_tree=None, sigma=None, X_enc=None):
9         g = loss.gradient(y, F_pred)
10        h = loss.hessian(y, F_pred)
11
12        # [*1] Langevin noise: epsilon ~ N(0, sqrt(2*lr/beta))
13        sigma_noise = np.sqrt(2.0 * lr / self.beta)
14        g = g + self.rng.normal(0.0, sigma_noise, size=g.shape)
15        return g, h
16
17    def update(self, F_pred, tree_pred, lr):
18        # [*3] shrinkage: F_t = (1 - alpha*lr)*F_{t-1} + lr*h_t
19        return (1.0 - self.alpha * lr) * F_pred + lr * tree_pred

```

7.5 Epistemic Uncertainty via Posterior Sampling

A single SGLB run can be interpreted as producing an approximate posterior-like sample. To estimate **epistemic uncertainty** (heuristically), run S independent chains (different random seeds) and compute the variance of predictions:

$$\hat{\mu}(x) = \frac{1}{S} \sum_{s=1}^S F^{(s)}(x), \quad \hat{\sigma}^2(x) = \frac{1}{S} \sum_{s=1}^S (F^{(s)}(x) - \hat{\mu}(x))^2$$

$\hat{\sigma}^2(x)$ is large when the S models disagree — indicating that the data does not strongly constrain the prediction at x . This variance is used as a **heuristic proxy for epistemic uncertainty**: the intuition is that it should decrease as more data constrains the model, unlike aleatoric uncertainty (irreducible data noise), though the connection to true Bayesian posterior uncertainty is approximate.

Listing 13 Epistemic uncertainty via ensemble (`model.py`)

```
1 # SGLBEnsemble acts as a factory only (boosting.py)
2 class SGLBEnsemble:
3     def make_sglb(self, seed) -> SGLB:
4         return SGLB(alpha=self.alpha, beta=self.beta,
5                     rng=np.random.default_rng(seed))
6
7 # GradientBoosting trains S independent models, stores tree_sets (model.py)
8 # fit():
9 for seed in self.boosting.seeds:
10     sglb = self.boosting.make_sglb(seed)
11     trees = self._fit_single(X, y, sglb, rng=np.random.default_rng(seed))
12     self.tree_sets.append(trees)
13
14 # predict_uncertainty():
15 preds = np.stack(
16     [self._predict_single(X_enc, trees, self.boosting.make_sglb(s))
17      for trees, s in zip(self.tree_sets, self.boosting.seeds)],
18     axis=0,
19 ) # (S, N)
20 return preds.var(axis=0) # epistemic uncertainty
```

7.6 Practical Considerations

Comparison with Random Forest variance. RF uncertainty is also estimated via ensemble variance, but the source of diversity is bootstrap sampling and feature subsampling. SGLB diversity comes from Langevin noise in gradient space. Both share the same structural limitation for tree-based models: extrapolation regions (out-of-distribution x) collapse to the nearest leaf, so all models agree and variance drops to zero. Within the training distribution, SGLB variance and RF variance capture similar epistemic uncertainty and tend to be highly correlated.

Comparison with aleatoric uncertainty (RMSEWithUncertainty). RMSEWithUncertainty trains a single model to jointly predict $\mu(x)$ and $\sigma^2(x)$ by minimizing the Gaussian negative log-likelihood: $L = \frac{(y-\mu)^2}{2\sigma^2} + \frac{1}{2} \log \sigma^2$. This $\sigma^2(x)$ measures *data noise* — it does not decrease with more data. SGLB $\hat{\sigma}^2(x)$ measures *model uncertainty* and does decrease with more data.

Cost. S independent SGLB runs multiply training time by S . For monthly retraining with moderate N , $S = 10$ – 20 is practical.

When to use SGLB instead of Ordered Boosting. For datasets with only numerical features and an existing walk-forward validation structure, Plain mode + SGLB provides regularization (via shrinkage and noise) and uncertainty estimation without the computational overhead of Ordered Boosting.

8 Module Architecture

The implementation is split into five modules, each with a single responsibility:

Module	Class / Function	Responsibility
tree.py	ObliviousTree	Split search, bit-index prediction
loss.py	MSE, LogLoss	Gradient, Hessian, leaf value
ordered_ts.py	OrderedTS	Categorical feature encoding
ordered_boost.py	ordered_gradients_from_leaf_ids	Per-leaf gradient estimation
model.py	CatBoostScratch	Assembly only — no logic

model.py sequences the modules in the correct order and holds no algorithmic logic of its own. Each module is independently testable.

Listing 14 Training loop in model.py (`_fit_single`, simplified)

```

1 def _fit_single(self, X, y, boosting_strategy, rng):
2     N, F_pred = len(y), np.full(len(y), self.F0)
3     trees, prev_tree = [], None
4
5     for _ in range(self.n_estimators):
6         sigma = rng.permutation(N)
7
8         # (a) Ordered TS: encode categorical features using permutation sigma
9         X_enc = self._ts_enc.transform(X, y, sigma) if self.cat_features else X
10
11        # (b) gradient + hessian via boosting strategy
12        g, h = boosting_strategy.compute_gradients(
13            y, F_pred, self.loss, self.learning_rate,
14            prev_tree=prev_tree, sigma=sigma, X_enc=X_enc,
15        )
16
17        # (c) fit tree on permuted samples
18        tree = self.tree_cls(max_depth=self.max_depth,
19                             min_samples_leaf=self.min_samples_leaf,
20                             ).fit(X_enc[sigma], g[sigma], h[sigma])
21
22        # (d) update ensemble
23        F_pred = boosting_strategy.update(F_pred, tree.predict(X_enc),
24                                         self.learning_rate)
25
26        prev_tree = tree
27        trees.append(tree)
28    return trees

```

9 Experiments

9.1 Setup

We test on a synthetic regression problem:

$$y = x_0^2 + 0.5x_1 + \varepsilon, \quad x \sim \text{Uniform}(-3, 3)^2, \quad \varepsilon \sim \mathcal{N}(0, 0.09)$$

with $N = 400$ samples (train 300, test 100), 100 estimators, $\eta = 0.1$, depth = 4. The noise variance $\sigma^2 = 0.09$ is the irreducible error floor — no model can achieve lower test MSE in expectation. A baseline that predicts the training mean achieves $\text{MSE} \approx 7.7$.

All six combinations share the same modular `GradientBoosting` interface:

```

1 model = GradientBoosting(
2     tree=ObliviousTree, loss=MSE(), boosting=OrderedBoosting(), ...
3 )

```

9.2 Results

#	Configuration	Train MSE	Test MSE
1	DecisionTree + MSE + Vanilla	0.0186	0.1798
2	ObliviousTree + MSE + Vanilla	0.0878	0.2294
3	ObliviousTree + MSE + Ordered	0.0786	0.1718
4	ObliviousTree + LogLoss + Ordered	0.0110	0.0694
5	ObliviousTree + MSE + SGLB	0.1089	0.2540
6	ObliviousTree + MSE + SGLBEnsemble	0.1038	0.2562

Configuration 6 additionally reports epistemic uncertainty: train = 0.00258, test = 0.00371 (test 44% higher).

9.3 Analysis

#1 (DecisionTree + Vanilla): severe overfitting. Train MSE $0.0186 \ll 0.09$ (noise floor) — the asymmetric tree memorizes the training data almost perfectly. Test MSE 0.1798 is twice the noise floor, suggesting that the flexibility that enables tight training fit hurts generalization.

#2 vs #1 (ObliviousTree): lower expressiveness, not always better. Symmetric splits constrain each depth level to a single condition, raising train MSE to 0.0878 (closer to the noise floor). However, test MSE 0.2294 is *worse* than #1. With only 100 estimators, the Oblivious Tree has not converged — its expressiveness limit acts more like underfitting than regularization. With more iterations this ordering may change, but should be tested.

#2 → #3 (Ordered Boosting): 25% test MSE reduction. Train MSE changes slightly ($0.0878 \rightarrow 0.0786$), while test MSE drops from 0.2294 to 0.1718 — a 25% improvement. This is the core effect of Ordered Boosting: the per-leaf cumulative correction removes the portion of each gradient already “explained” by preceding samples, using prefix-corrected gradient estimates that reduce self-influence and may generalize better.

#4 (LogLoss + Ordered): classification baseline. The target is binarized at the median, giving a balanced 50/50 split. A random classifier achieves $\text{MSE} = 0.25$; the model achieves 0.0694 on the test set, suggesting effective learning. The train/test ratio ($\approx 6\times$) is larger than in regression, partly because the 0-1 loss surface is non-convex and the model fits the training labels tightly.

#5–6 (SGLB): regularization at a cost. With $\beta = 10$ ($\sigma_\epsilon = \sqrt{2\eta/\beta} = 0.141$), SGLB raises train MSE to 0.1089 — the noise injection prevents the model from fitting the training set tightly, as intended. However, test MSE 0.2540 remains worse than Ordered Boosting (0.1718).

The reason is structural: SGLB’s benefit (escaping local minima, stochastic regularization) requires either a multimodal loss landscape or large N . The synthetic target $y = x_0^2 + 0.5x_1 + \epsilon$ is nearly convex, so there are no local minima to escape. Ordered Boosting removes a provable bias; SGLB adds beneficial noise whose advantage depends on the problem geometry.

SGLBEnsemble (#6) averages 10 independent chains, marginally reducing variance without changing the bias, giving results nearly identical to a single SGLB run (#5).

Epistemic uncertainty (#6): correct ordering. Test uncertainty (0.00371) exceeds train uncertainty (0.00258) by 44%. Test samples lie further from the training distribution’s center, causing the 10 Langevin chains to land in different leaves more often — producing higher inter-model variance. This is consistent with the ensemble variance serving as a proxy for *epistemic* uncertainty: regions less covered by training data yield higher prediction disagreement.

9.4 Summary

Observation	Implication
#1 overfits, #2 underconverges	Tree choice interacts with iteration budget; symmetric trees need more estimators.
#2 $\rightarrow$ #3 (−25% test MSE)	Ordered Boosting removes a provable gradient bias; benefit is consistent.
#5–6 worse than #3	SGLB regularization is problem-geometry dependent; use when N is large or loss is non-convex.
Uncertainty: test > train	SGLB Ensemble variance is consistent with epistemic uncertainty interpretation.

10 Limitations and What We Left Out

This implementation covers the core ideas but omits several features present in production CatBoost:

- **Multi-permutation variance reduction.** The original paper uses s independent permutations per iteration and averages the resulting gradients. This reduces the variance introduced by the single-permutation approximation.
- **GPU-optimized split search.** CatBoost’s GPU kernel exploits the Oblivious Tree structure for fully parallel split evaluation.
- **Advanced categorical handling.** One-hot encoding for low-cardinality features, frequency encoding, and combinations.
- **Early stopping and overfitting detector.**
- **Multiclass and ranking losses.**
- **Stochastic Gradient Langevin Boosting (SGLB) and posterior sampling.** The `posterior_sampling=True` option in CatBoost activates SGLB (Ustimenko & Prokhorenkova, 2021), which modifies the boosting update as:

$$F_t = (1 - \alpha\eta) F_{t-1} + \eta \cdot h_t^{\text{noisy}}, \quad g_k^t \leftarrow g_k^t + \epsilon_k, \quad \epsilon_k \sim \mathcal{N}(0, \sigma^2)$$

The shrinkage term $(1 - \alpha\eta)$ regularizes the model; the noise ϵ_k injected into each gradient causes the Markov chain of models to converge not to a point estimate but to a distribution $p(F) \propto \exp(-L(F)/\eta)$, which approximates the posterior over functions under a uniform prior (in the continuous-time, infinite-data limit). Running multiple independent chains and averaging their predictions yields an approximate Monte-Carlo estimate of posterior variance — a useful but heuristic uncertainty estimate. This is particularly useful when Plain mode (no Ordered Boosting) is used with numerical features and an external walk-forward validation structure already handles look-ahead bias, as the SGLB noise provides regularization without the computational overhead of Ordered Boosting.

11 Conclusion

Starting from undergraduate calculus and the idea of gradient descent, we derived and implemented five innovations that transform Vanilla GBM into a modern, principled gradient boosting system:

1. **Newton Step** generalizes the leaf value computation via a second-order Taylor approximation, enabling any differentiable loss function through a unified gradient/Hessian interface.
2. **Oblivious Trees** replace standard decision trees with a symmetric structure that enables $O(1)$ bit-index prediction and efficient GPU parallelism.
3. **Ordered Target Statistics** prevent categorical feature encoding from leaking the target by conditioning on strictly preceding samples in a random permutation.
4. **Ordered Boosting** computes CatBoost-style prefix-corrected gradient estimates by computing each sample’s residual against a model that has never seen that sample, implemented efficiently via per-leaf cumulative statistics.
5. **Posterior Sampling (SGLB)** injects Langevin noise into gradient estimates and applies shrinkage, transforming the boosting Markov chain into an approximate posterior sampler over functions. Running S independent chains yields heuristic epistemic uncertainty estimates via Monte-Carlo variance.

Each innovation is independently testable and motivated by a concrete limitation of the previous approach. The experimental results are consistent with the theoretical predictions: Ordered Boosting reduces test MSE by reducing gradient self-influence, and SGLB uncertainty estimates assign higher variance to out-of-distribution regions within the training support.

References

- [1] Prokhorenkova, L., Gusev, G., Vorobev, A., Dorogush, A. V., & Gulin, A. (2018). *CatBoost: unbiased boosting with categorical features*. Advances in Neural Information Processing Systems (NeurIPS), 31.
- [2] Friedman, J. H. (2001). *Greedy function approximation: a gradient boosting machine*. Annals of Statistics, 29(5), 1189–1232.